

## Assignment: Human Keypoint Detection in an Image

### Objective:

Build a Python script that accepts a single input image and uses a pose estimation library to determine whether a human is present in the image.

### Requirements:

- Use **Mediapipe** for pose/keypoint detection (we recommend this for its simplicity and strong performance).
- If keypoints are detected, output:
  - A message such as “**Human detected**”.
  - Optionally, the number of keypoints or their coordinates.
  - You can also overlay small dots on each detected keypoint and save the image to disk (optional but appreciated).
- If no human is detected, simply output: “**No human detected.**”

You do **not** need to use webcam input—just work with one uploaded image.

### Tools:

- Python (any version you're comfortable with).
- OpenCV and Mediapipe (e.g., `pip install opencv-python mediapipe`).

### Estimated Time to Complete:

According to ChatGPT, this assignment should take approximately **2–3 hours** to complete, assuming a working knowledge of Python and the libraries mentioned above.